

Research Project

Statistical Foundations of Machine Learning

Ward Gauderis & Fabian Denoodt

27/06/2022

Vrije Universiteit Brussel

Summary

1. Stochastic noise and weight decay regularisation
2. Support vector machine kernel comparison
3. Decision tree and k-nearest neighbours regressor forecasting

Stochastic noise and weight decay regularisation

What is the influence of stochastic noise in the dataset on the in- and out-of-sample error of a neural network and how does weight decay regularisation counter this?



The datasets

Synthetic normalised (x, y) with $x \in \mathbb{R}^2, y \in \{0, 1\}$

Controllable stochastic noise:

- **Data noise:** $x' = x + \epsilon$ with $\epsilon \sim \mathcal{N}(0, \sigma^2)$
- **Label noise:** swap fraction α of labels y

No confounding deterministic noise



The models

Neural network

- Predict $h(x) = P(y = 1|x)$
- Linear transformations with non-linear differentiable ReLU activation functions
- Maximise likelihood by minimising the cross-entropy error
- (20, 20) hidden layers & 2000 epochs

Augmented error regularisation:

$$E_{aug}(h, \lambda, \Omega) = E_{in}(h) + \frac{\lambda}{N} \Omega(h)$$

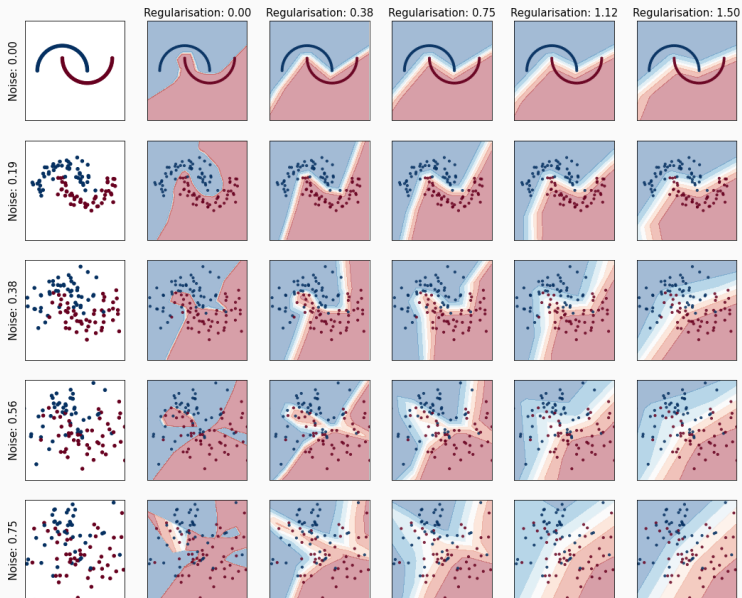
Weight-decay with L_2 norm:

$$E_{aug}(w) = E_{in}(w) + \frac{\lambda}{N} ||w||^2$$




Experimental setup

Repeat for every combination of dataset type, label and data noise:


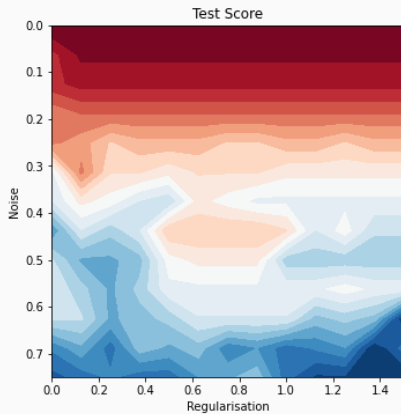

1. Generate 13 datasets of size 100 with noise $\in [0, 0.75]$
2. Create 13 models of with regularisation $\in [0, 1.5]$
3. Train every model on every dataset
4. Compare decision boundaries and training and testing accuracies




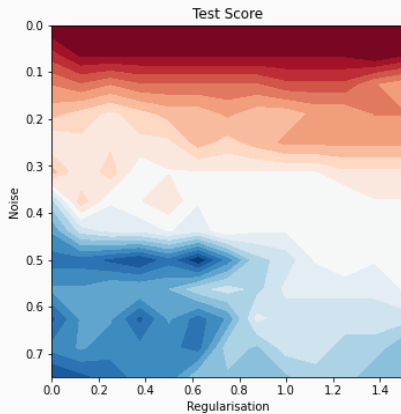

Moon dataset with data noise



Moon dataset with data noise



Circles dataset with label noise



Moon dataset with combined noise

Conclusions

- Neural networks with less effective parameters generalise better on noisier datasets
- Label noise versus data noise
 - more detrimental to generalisability
 - can be combated with less regularisation
- Underfitting is less punishing than overfitting
- E_{in} becomes less informative about E_{out} with increasing noise and regularisation
- Optimal λ is hard to know up front without data snooping and should be chosen through model selection

Support vector machine kernel comparison

How do the linear kernel, polynomial kernel and radial basis function compare to each other, when applied to a synthetic two-dimensional dataset?

Decision tree and k-nearest neighbours regressor forecasting

How does the decision tree regressor model compare to the k-nearest neighbour regressor model in terms of in- and out-of-sample error for time series forecasting?

Experimental setup

Root-mean-square error:

$$E = \sqrt{\frac{\sum_{n=0}^N (y_n - h(x_n))^2}{N}}$$

Regression score:

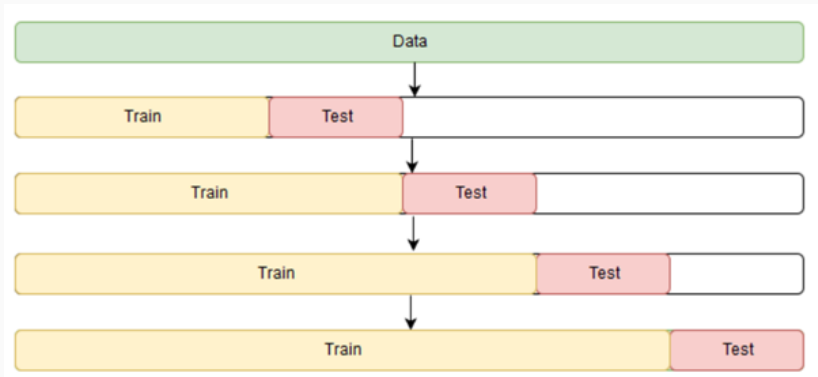
$$1 - \frac{\sum_{n=0}^N (y_n - h(x_n))^2}{\sum_{n=0}^N (y_n - \bar{y})^2}$$

Data:


1. Split into subsequent training and testing set
2. Process normalised data

Model comparison:

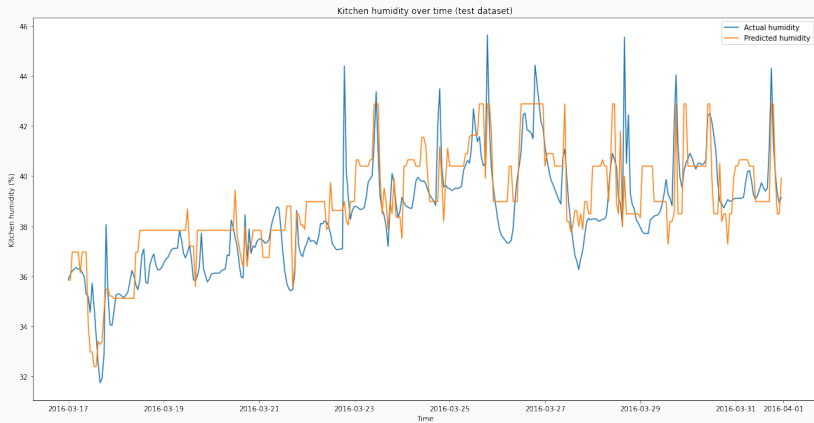
1. Baseline models
2. Equally tuned models through randomised grid search



Time series cross-validation for forecasting




Baseline decision tree regressor




Baseline decision tree regressor

	Training	Testing
Score	1.000	0.479
RMSE	0.000	1.551



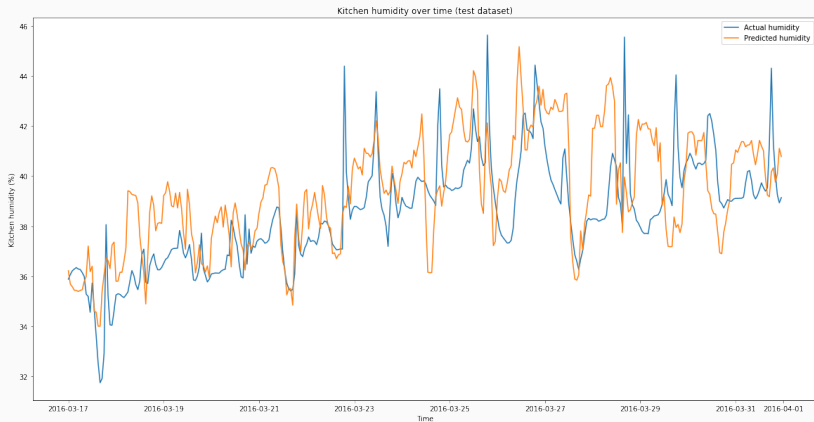
Tuned decision tree regressor

	Training	Testing
Score	0.949	0.554
RMSE	0.871	1.434



Baseline k-nearest neighbours regressor

	Training	Testing
Score	0.945	0.040
RMSE	0.908	2.105



Tuned k-nearest neighbours regressor

	Training	Testing
Score	1.000	0.089
RMSE	0.000	2.050

Conclusions

- Tuned decision tree regressor is most likely to generalise best
- Model nature is visible in predictions:
 - Decision tree predicts conservative smooth surfaces
 - K-nearest neighbours predicts noisy erratic changes resembling the training data
- Time-series forecasting is extrapolation
- Model selection can improve E_{out} by reducing or increasing model complexity
- Data preprocessing and feature selection is important and can be guided also by cross-validation